



# OpenFlightHPC Bootcamp

*Liverpool HPC MSc Program*

*Next Generation HPC*

February, 2020

# Agenda



12:00 Introductions

*What is Next-Generation High Performance Computing?*

*The HPC Market*

*Cloud HPC: A (sort of short) History*



13:00 OpenFlightHPC project overview

13:30 Exercises in OpenFlightHPC

*HPC cluster applications and workloads*

*Next-generation workloads*

*Review, wrap-up and questions*



16:00 Close





# Next Generation HPC

*Alces Flight*

February, 2020

# Alces group – what is it?



- UK small/medium sized enterprise (SME)
- Private self-funded company
- Specialist in High Performance Computing
- Formed in 2008
  - >100% YOY growth for FY1-5
  - No external funding or shareholders
  - Run of 42 profitable quarters
  - Parent + trading divisions



# Alces group – what does it do?



- Builds Research Computing environments for customers
- Acts as a service provider
  - Logistics team physically delivers and installs equipment
  - Engineering team manages and monitors customer systems
  - Software teams responsible for:
    - Building the tools and products used to deliver our services
    - Creating new products to help solve customer problems
    - Supporting our consultants working with 3<sup>rd</sup>-party software

# Alces group – what does it do?



- Project examples:
  - Configuring a bare-metal HPC cluster for a University
    - e.g. Barkla HPC cluster at Liverpool
  - Extending an existing bare-metal facility with new technologies
  - Reconfiguring an existing facility to do a new thing
  - Advising customers on creating a new HPC facility using hardware, public or private cloud resources
  - Enabling on-demand research computing environments for dynamic end-user requirements

# Alces group – what does it do?



- Technology examples:
  - Hardware
    - IBM, HPE, Dell, Cray, Supermicro, Intel, Mellanox, Cisco....
  - Cloud
    - Amazon Web Services, Microsoft Azure, Google Compute cloud
  - HPC software
    - Linux, job-scheduler, Docker, CUDA, MPI, compilers, languages
  - Application software
    - NAMD, Gromacs, Tensorflow, Lightwave, WRF, Code Saturne, MapReduce, Jupyter Notebook, Matlab

# Team members

- Technical Sales Consultant
  - Wil Mayers
- Change Management lead
  - Cristin Merritt
- Technical Product Manager
  - Stu Franks



# Traditional High Performance Computing



- Serial job
  - One CPU core / node / task at a time
  - If dataset is infinitely divisible these are embarrassingly parallel
  - Often run in workflows or batches (e.g. batch-cluster)
- Multi-threaded jobs
  - Like serial jobs, but use multiple resources on the same node
- Parallel jobs
  - Use resources (CPU/memory) from multiple nodes
  - Commonly use an API to simplify programming
  - Message passing interconnect (MPI) handles communication
  - Scalability depends on both application and system



# Speciality HPC workloads

- Distributed / big-data jobs
  - Lots (and lots and lots) of data
  - You want to do lots of work on the same dataset
  - Compute is more efficient to run where the data is already located
  - Designed for large scale workloads
    - Fault-tolerance built in to the software framework
    - Difficult to get smaller workloads to run efficiently
- Artificial intelligence / machine-learning
  - Training a computer model to make statistical predictions
  - Outputs typically include a “confidence” rating
  - Works well on many-core processors (e.g. GPUs)
  - An example of an “overlay” technology
    - Can help improve/optimise other workflows

# High Performance Workloads



- Sometimes just one application
  - Matlab, Python code, R script, Java script
  - Gene sequence searching, mapping, detection
  - NAMD, Gromacs, Quantum Espresso
- Sometimes a workflow of applications
  - Prepare data -> process data -> recombine outputs -> present
  - Sieve data -> compare found results -> present
  - Use Hadoop to find pictures of animals on the web -> Use AI to detect which animals are cats -> Use python array to organise by geographic region -> How many Liverpool residents have cats?



# Where do these applications come from?



- Home-grown
  - Written by you, your lab partner, your tutor
- Open-source
  - Written and published by a group of developers
  - Mixture of small, single-source and massive group efforts
  - Myriad of software licenses, EULAs
- Commercial
  - Can be closed code-base, proprietary
  - Could be optimised/supported/tested version of open-source app
  - You might need to pay to use it

# Distribution of HPC software apps

- Compile it from source-code
  - Pretty standard for open-source HPC cluster applications
- Download binary packages
  - Download, install and run a Windows app
  - RPM, DEB, DMG packages for Linux/Mac
- Get it in a container
  - Docker, Singularity, VMware, Kubernetes, Virtualbox VM
  - AWS AMI, Azure or GCP disk image
  - Often includes app, libraries and supporting bits of operating system
- Launch or use an existing online service; e.g.
  - Github / Gitlab
  - Jupyter Notebook
  - Tensorflow.org

# Workload maturation



- Most work begins in “discovery” mode
  - Try something, see if it works, try something different
  - Rapid prototyping, agile software development
  - Fast, innovative, quickly changing
- Compute resources tend to be
  - Low capacity, with limited scalability testing
  - Leveraging third-party software, often from multiple sources
  - Very variable during the development process
  - Unoptimised, inefficient, difficult to scale

# Workload maturation



- As workload matures
  - We have better idea of what compute we need
  - Centralise on a specific software environment
  - We can begin to optimise for scale, performance, cost
- Compute resources tend to be
  - Less generic and more specific
  - Capable of running unattended
  - Scaled to handle bursts of work
  - Cost optimised

# Workload lifecycle



- High Performance Computing is all of these things
  - Unoptimised, adhoc, rapidly changing
    - Launching an AWS AMI created by a vendor to test a theory
    - Using a vendor-supplied service to test something
    - Trying your MapReduce code on a public Hadoop service
    - Testing your gene sequence for HG19 matches on Galaxy.org
  - Tightly-coupled, highly optimised, reproducible
    - Running 100 x NAMD chemistry simulations on Barkla HPC cluster
    - Submitting 2,000 tensorflow simulations to a batch cluster
    - Connecting to the pay-per-use Azure MapReduce service to process data transmitted from self-driving cars



# What is Next-generation HPC?



- All those applications...
  - ...run as any of those workflows
  - ...using software from any of those sources
  - ...deployed using any of those mechanisms
  - ...on any of those platforms
  - ...by end-users of mixed abilities
  - ...over the workload's entire lifecycle.
- Plus the things that were invented while we just read that.



# The HPC Market

1960-2020

February, 2020

# Traditional HPC cluster market



- Most traditional HPC users run on HPC clusters
  - Central service shared between many users
  - Lots of compute nodes, each running separate OS
  - Shared storage
  - Fixed software environment with job-scheduler for time-sharing
  - One cluster typically lasts 3-5 years then is replaced
- UK has lots of these
  - National facility is called “Archer”, hosted at EPCC
  - Liverpool University has “Barkla” HPC cluster
    - 50-100 similar machines across the UK

# On-prem HPC market limitations



- Only so much HPC work to do
  - Large capital investment required = high barrier to entry
  - HPC applications somewhat specialised to setup + use
- Logistical requirements cause limitations
  - Tricky to deliver and install new clusters outside local area
  - HPC cluster installations are often loss-making for vendors
- Lack of standardisation and repeatability
  - Every installation is it's own special snowflake
  - Irreproducibility reduces value of research outputs

# Addressing a changing market



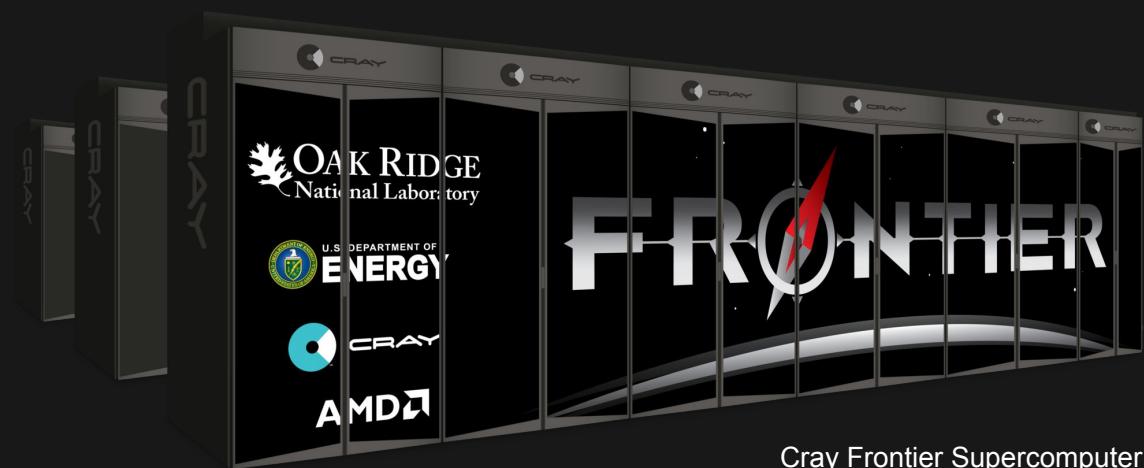
- *Only so much HPC work to do*
  - Make HPC more approachable
  - Reduce barrier to entry
- *Logistical requirements cause limitations*
  - Stop relying on hardware
  - Build a fairer, sustainable market
- *Lack of standardisation and repeatability*
  - Standardise
  - Educate

# Where HPC fits into the IT Industry

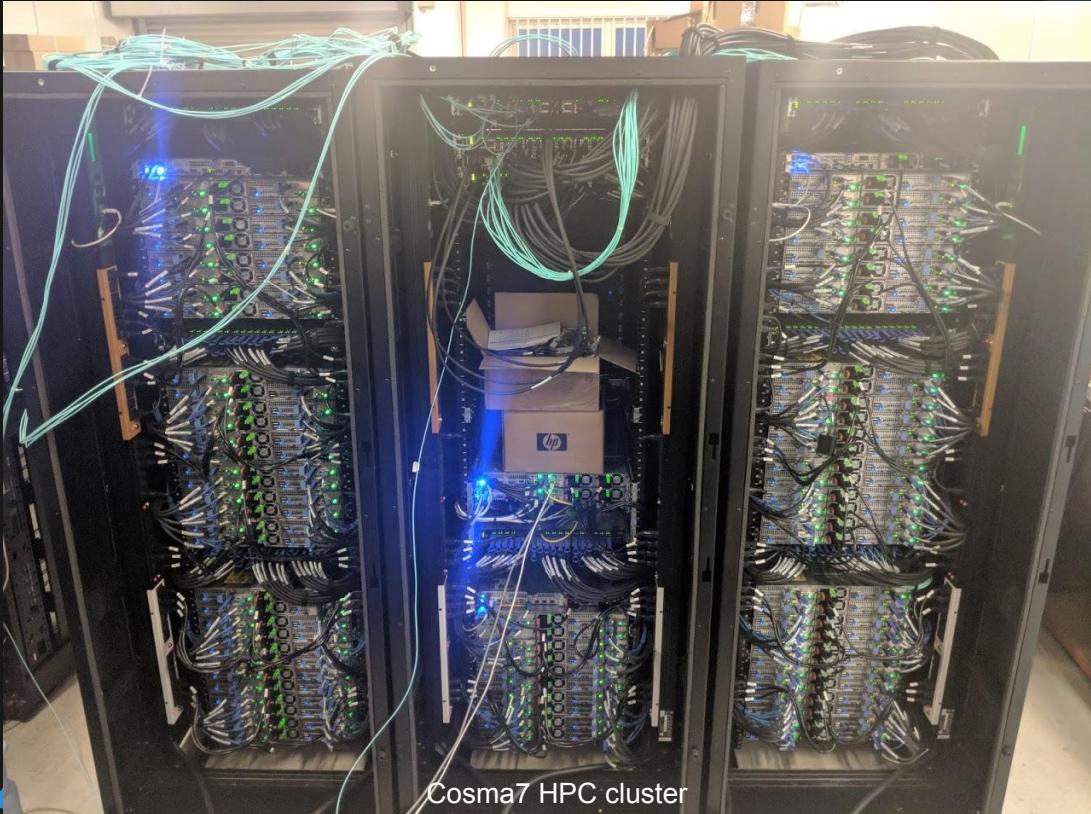


Ferrari 812 Superfast

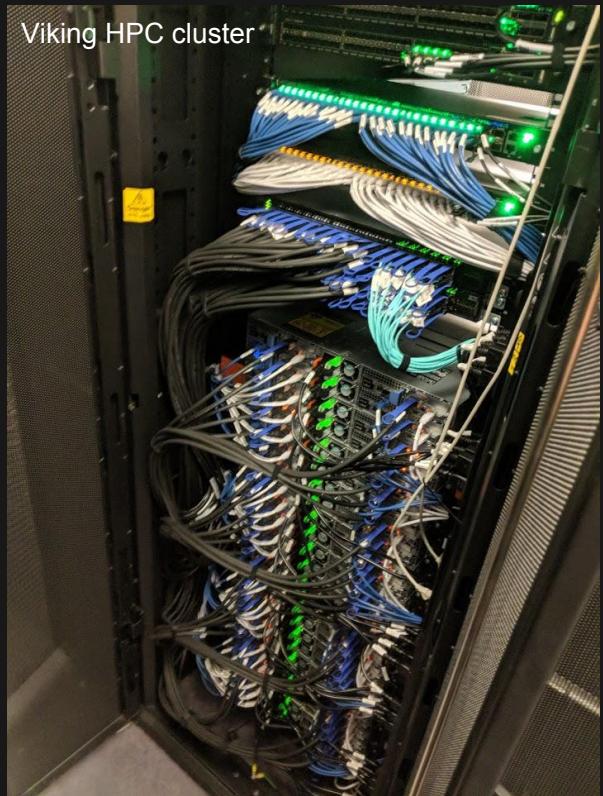
# Where HPC fits into the IT Industry



# Where HPC fits into the IT Industry



Viking HPC cluster



# Where HPC fits into the IT Industry



Ferrari 812 Superfast



Landrover Defender TD5



Vauxhall Corsa 1.2



Trex Fly 5 bicycle



30 seater school bus

# Cost, not technology, is now the key factor



- Commodity computers and storage
  - Allow low-cost HPC clusters to be created
  - Very good for well-developed workloads
  - Inflexible for discovery workloads
- The rise of cloud computing
  - Much easier to perform discovery work
  - On-demand launch, try it out, throw it away, try again
  - Much lower barrier to entry with access to latest technology
  - Currently more expensive for static workloads...

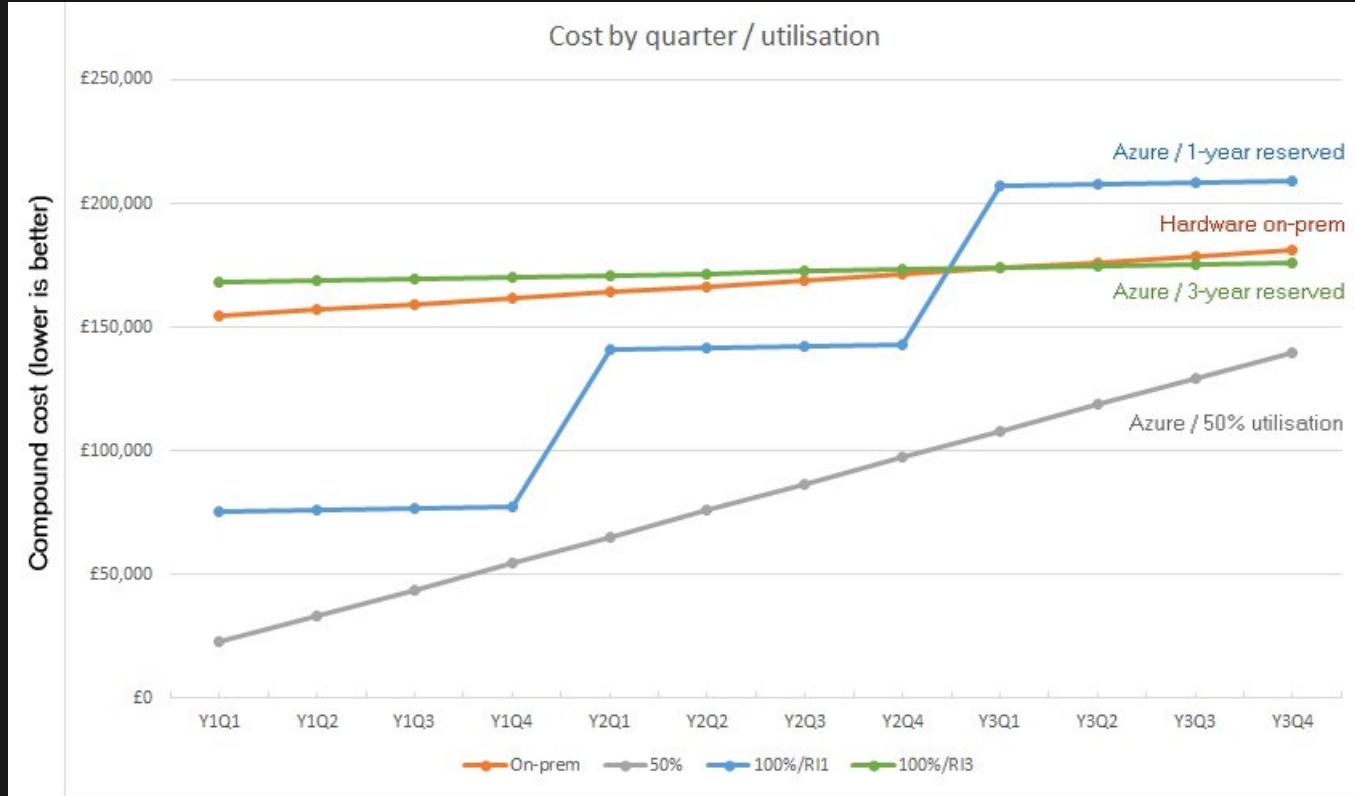
# Is Cloud more expensive than on-prem?



- Buying and running your own hardware
  - Simple to manage your budget
  - Difficult to quantify + measure all the real costs
  - Multi-year investments in buildings, people, process, etc.
  - Measure cost-effectiveness in arrears
- Cloud computing
  - On-demand availability vs. committed spend to get discounts
  - Few demands on local infrastructure; great for start-ups
  - Sometimes hard to predict the final bill
  - Measure cost-effectiveness day-by-day



# On-prem vs Cloud examples



# On-prem vs cloud examples



- NHS HPC cluster
  - Gene therapy unit, Manchester Children's Hospital
  - Cluster used to diagnose and treat patients
  - Samples taken, rapidly sequenced and processed
  - Patients arrive at 9am and leave by 5pm with results
- On-prem HPC cluster can process one child per day
  - 22% utilisation (9-5, Mon-Fri, 52 weeks p/a)
- Cloud HPC cluster grows/shrinks on-demand
  - 3x cheaper to run (or 3 children per day for same cost)

# HPC market futures



- IDC predicting migration to cloud
  - Some challenges for very specialist HPC machines
  - Cloud technology already good enough for most small-medium HPC clusters
- As cloud takes over IT, hardware becomes more expensive
  - Owning your own data-centre becomes rarer
  - Hardware vendors downsize as economy of scale is reversed
  - Building an on-prem HPC cluster becomes more expensive
- Skills-gap in building, deploying and managing workloads



# *Cloud HPC: A (sort of short) History*



*Cristin Merritt*

*Alces Flight Cloud Adoption Project*

February, 2020

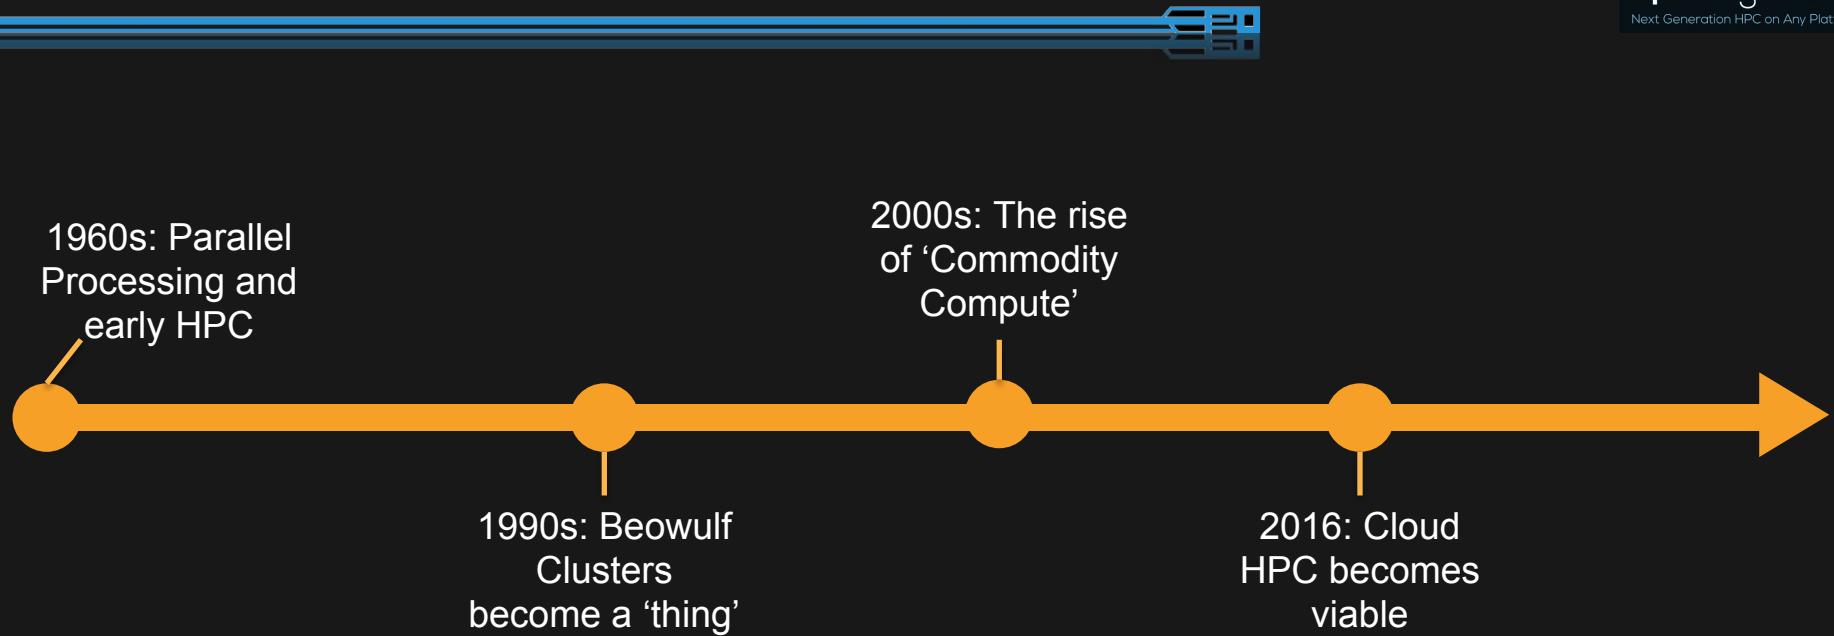
# Alces Flight at Liverpool University



- Integrators for Liverpool since 2017 (Barkla Cluster)
- Provide Managed Services for Liverpool
  - Look after the daily checks, general maintenance, and manage/perform change requests
  - Build cloud HPC projects (AWS and Azure for Liverpool)
  - Allow on-site team at Advanced Computing Facilities to build user expansion
- Launched OpenFlightHPC Project in November, 2019...

*But a bunch of stuff happened before any of this...*

# A poorly drawn HPC Timeline



# 2016: Cloud HPC (versus Hardware)



- Launch of Alces Flight Compute cluster product
  - A traditional HPC compute cluster running on public cloud
- Launch of Cloud Adoption Project for HPC
  - Tracked subscriptions on AWS and Azure (total of around 275 sites... 222 of those are still considered active)
  - Five sites agreed to complete profile of use

*We learned a lot of things...*



# Cloud HPC disrupts HPC culture



- You cannot see/hold the hardware
- Technically anyone can do HPC on the cloud (purchasing barrier quite low)
- Practically no barriers to having the ‘newest kit’ running your workload
- Cloud opens the door for a lot of new fields with ideas on how to leverage HPC (skills shortage and gap emergence)
- Cloud allows projects to run as fast or as slow as you desire (time v. cost / workload/job queuing)

# Good things about cloud for HPC



- It does really well when it's project-based
  - Much easier to collaborate with others
- It loves open-source
  - Use whatever software you can find
- It does very well when work is embarrassingly parallel
- It is a great tool to make longer-term decisions with
  - Types of hardware to purchase
  - Testing new theories and ideas
  - Taking ideas into production



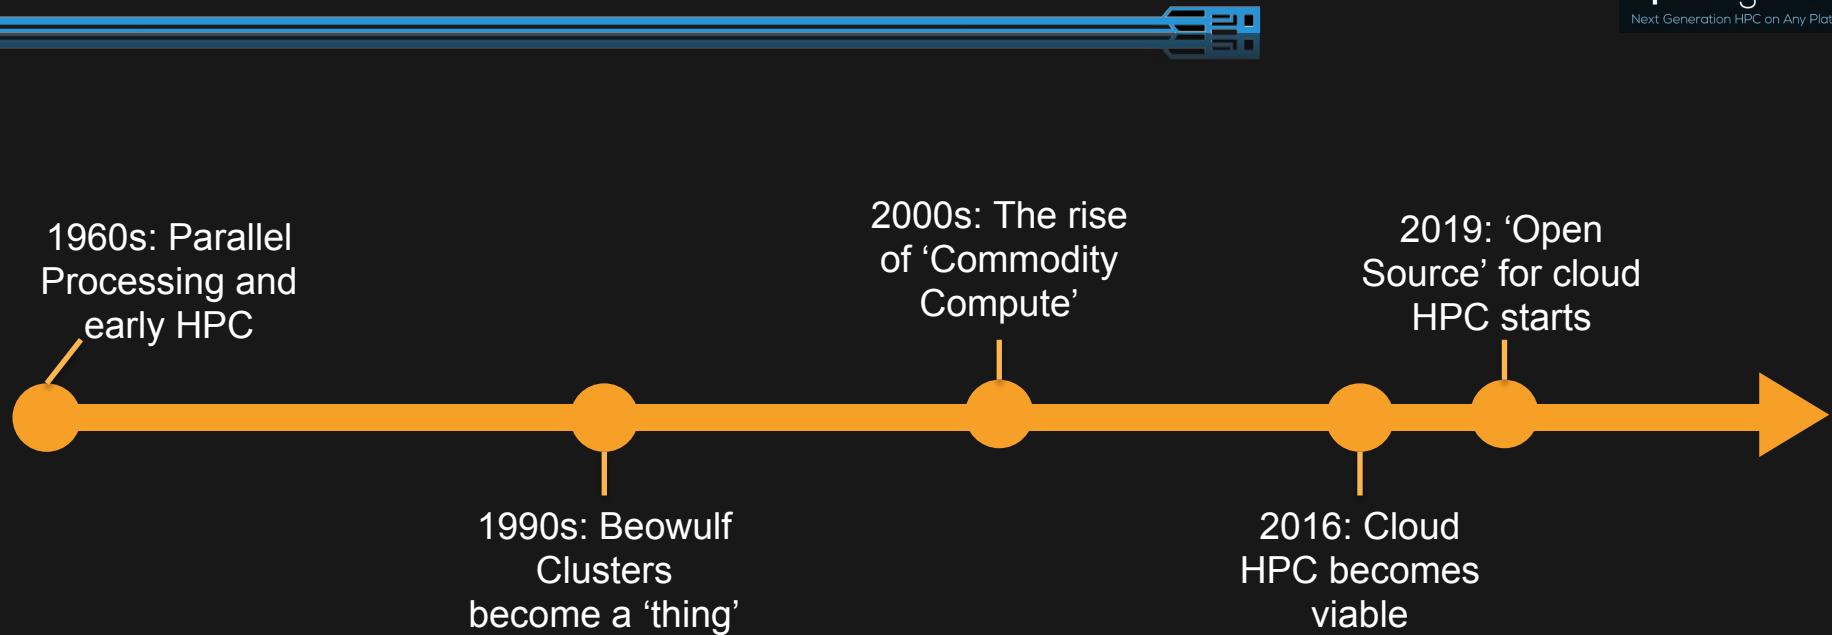
# Bad things about cloud for HPC



- Despite what the cloud vendors tell you, cloud HPC can get very expensive
  - Budget is not fixed, so you need to keep control
- Commercial vendors still don't quite know what to do about cloud
- Over abundance of value-add 'services'
  - Complex environment to navigate for a beginner
- Many parts means reproducibility can be an issue
  - How did I get here?



# A poorly drawn HPC Timeline

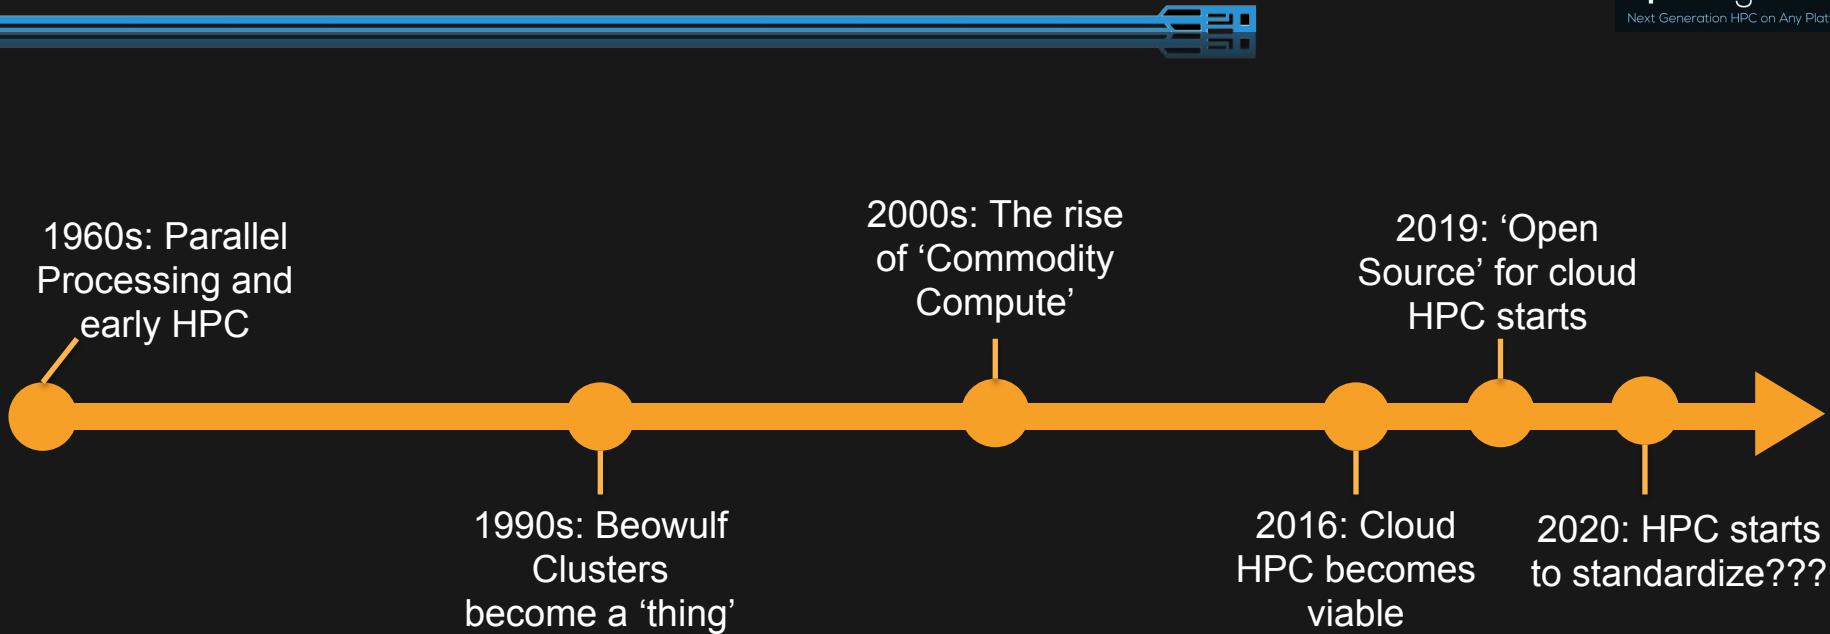


# 2019: Open-source for cloud HPC



- Subscription services to cloud HPC deemed too complex to maintain.
- Support for cloud HPC is difficult for cloud vendors to provide – especially as some level of customization work seems to happen in HPC projects.
- Open-source far more viable as you can optimize 75-90% across ANY platform.

# A poorly drawn HPC Timeline



# 2020: Real Skills



- People skills are ‘in’
  - Customer/Client/User-facing roles
  - Working in partnership or collaboration
- Niche skills are in – especially in Artificial Intelligence or Machine Learning optimization
- Services creation is in – see ‘people skills’ above.
- Standards are in... it’s not the fastest hardware it’s getting things done consistently and are easy to replicate.

# 2020: Where the jobs are...

- Research Software Engineers (RSEs): Client-facing roles dedicated to getting workloads to run well. (Field still new – launched approx. 2014)
- Software/Solution Architects (SAs): Commercial version of RSEs, often working with specific clients to build/optimize workloads.
- Computational Data Scientist: Focus on optimizing the access and use of large data sets – often deals in storage. (Might see ‘Storage Engineer’ in prospective job titles)
- CFD / Visualization Engineers: Strong on cloud development and proven to speed along projects which were once workstation based.
- Software Engineering: Generally focused in specific fields – demands in earth sciences (oil and gas) government (military, space exploration, weather) and open-source (ex. Redhat, SuSE, AWS, Azure...)

# 2020: Where Alces Flight is...



- Focused on standardization, automation and process of HPC services.
- Focused on utilizing open-source as a means of allowing our clients to work across platforms.

Launched OpenFlightHPC as a centralized project to help users to access, develop and run HPC workloads across any compute platform.



Next Generation HPC on Any Platform

## *Part One: Quick Introduction*

February, 2020

# Bootcamp agenda



- Part One: What is OpenFlightHPC?
  - The project goals
  - How you can get involved
- Part Two: Familiarisation exercises
  - Running applications on a cloud-based HPC cluster
  - Running a big-data workload
  - Running a machine-learning workload

# What is OpenFlightHPC?



- A *central resource* to help users run **HPC workloads** on any compute **platform**
- **HPC workloads**
  - Traditional batch and parallel HPC cluster jobs
  - Containerised / packaged applications
  - Client / server web services
- **Multi-platform**
  - Bare-metal
  - Private-cloud, public-cloud, online services

# What can OpenFlightHPC do for you?



- Help to provide you with what you need to run your HPC workload
- Connect you to other users running similar jobs
- Provide guidance for improving and optimising your work
  - Getting started with your discovery process
  - Performance optimisation
  - Improved accessibility
  - Process repeatability and sustainability
  - Budget control and cost reduction

# Example – image processing workload



- Initial discovery work
  - Performed using individual AWS instances
  - Iterative process improving successive images over time
- Scalability improvements
  - Host your existing image on larger / more-capable instances
  - Build a workflow that allows your job to run unattended
- Move to production
  - Use an autoscaling batch HPC cluster to process data on arrival
  - Run on the existing Barkla HPC cluster using spare cycles

# Example – MapReduce workload



- Initial discovery work
  - Run Apache Hadoop on a single machine
  - Perform test runs with different data to ensure your job works
- Scalability improvements
  - Upload your data to Amazon EMR service
  - Run larger queries to ensure your code works at scale
  - Consider an AI decomposition to narrow search parameters
- Move to production
  - Ensure your workflow is portable across a range of services
  - Review cost of using EMR vs. building your own Hadoop service

# How can you get involved?



- Use it
  - Bootcamp provides some example of workflows
  - Plan your future workloads, looking for the applications you want to use
  - Think about how you want to run your workload
    - Where will your software come from?
    - How will that software be distributed and deployed?
    - How can you run at scale?
    - Once things are running, how will you optimise your jobs?
- Share your experiences
  - Provide information for your users, stakeholders and researchers
  - Help other users on the OpenFlightHPC community site
  - Publish your own use-cases, success stories and research outcomes



Next Generation HPC on Any Platform

## *Part Two: Introductory Bootcamp*

February, 2020



# OpenFlightHPC familiarization



- You will need:
  - A laptop with a web-browser
  - Wireless internet access
- OpenFlightHPC web access:  
<http://bootcamp.openighthpc.org/>
- Alternatively:
  - An SSH client and a VNC client
    - Windows: <http://tiny.cc/turbovnc-win64>

# OpenFlightHPC familiarization



- What are you logging in to?
  - A compute cluster running in public cloud
    - One login node (*gateway1*)
    - Two small compute nodes (*node01* and *node02*)
    - Virtual private internal cluster network
    - Small, shared filesystem
    - Slurm job-scheduler
- What are we going to do with it?
  - Navigate the cluster, create and run some batch jobs
  - Install and configure some serial and parallel applications
  - Setup Apache Hadoop and run a MapReduce example
  - Run a job using an AI application container



# Logging in and cluster navigation

```
# id  
# pwd  
# df -h .  
# sudo pwd  
# sudo yum install dos2unix  
# ping node01
```



Press “y” to confirm

```
# flight start  
# ssh node01  
# df -h .  
# exit  
# sinfo -NI  
# flight info  
# flight env avail  
# flight set always on
```

# Linux HPC cluster application management



- Complexity of application varies
  - Some are simple; many are not
  - Most use a variety of different libraries
  - Some also use APIs (e.g. message-passing)
  - Version control is important for reproducibility
  - Dependencies must be satisfied across all nodes
- You can choose how to manage this
  - Configure things manually
  - Get help

# Compiling and running an app manually



```
# nano hello.c
```

Enter this content

```
#include <stdio.h>
int main() {
    printf("Hello, World!\n");
    return 0;
}
```

```
# gcc -o hello hello.c
# ls
# ./hello
Hello, World!
```

*Press CTRL+X to exit and enter “y” to save*



# Running a batch job on your cluster

```
# sinfo -NI  
# nano myjob.sh
```

Enter this content

```
#!/bin/bash  
#SBATCH -N1  
echo "Hello from $HOSTNAME"  
date  
sleep 5  
. /hello
```

```
# sbatch myjob.sh
```

Submitted batch job 2

```
# cat slurm-2.out
```

Press *CTRL+X* to exit and enter “y” to save

Use this number here

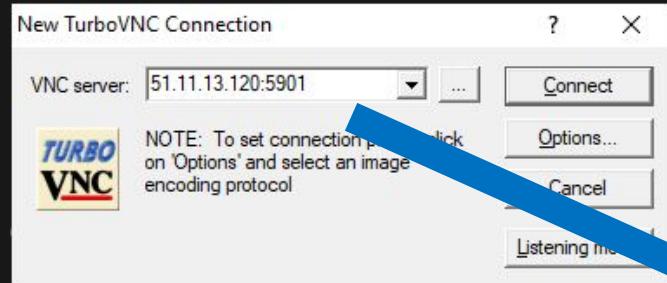
# Command-line vs graphical desktop

- Depending on your workflow, you might use
  - A command-line interface
  - A graphical desktop interface
  - A web-browser (locally or on the cluster)
- Graphical interaction is great for
  - Discovery work
  - Debugging an issue
  - Anything visual or interactive
- Command-line interfaces are good for
  - Unattended jobs
  - Anything that outputs a lot of data

# Starting a graphical desktop on a cluster

```
# flight desktop start gnome
```

*Start your local VNC client*



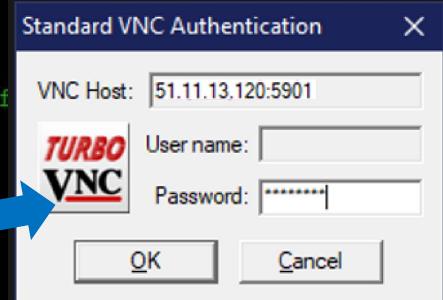
```
starting a 'gnome' desktop session:
```

```
>  Starting session
```

```
A 'gnome' desktop session has been started.
```

```
-- Session details ==
```

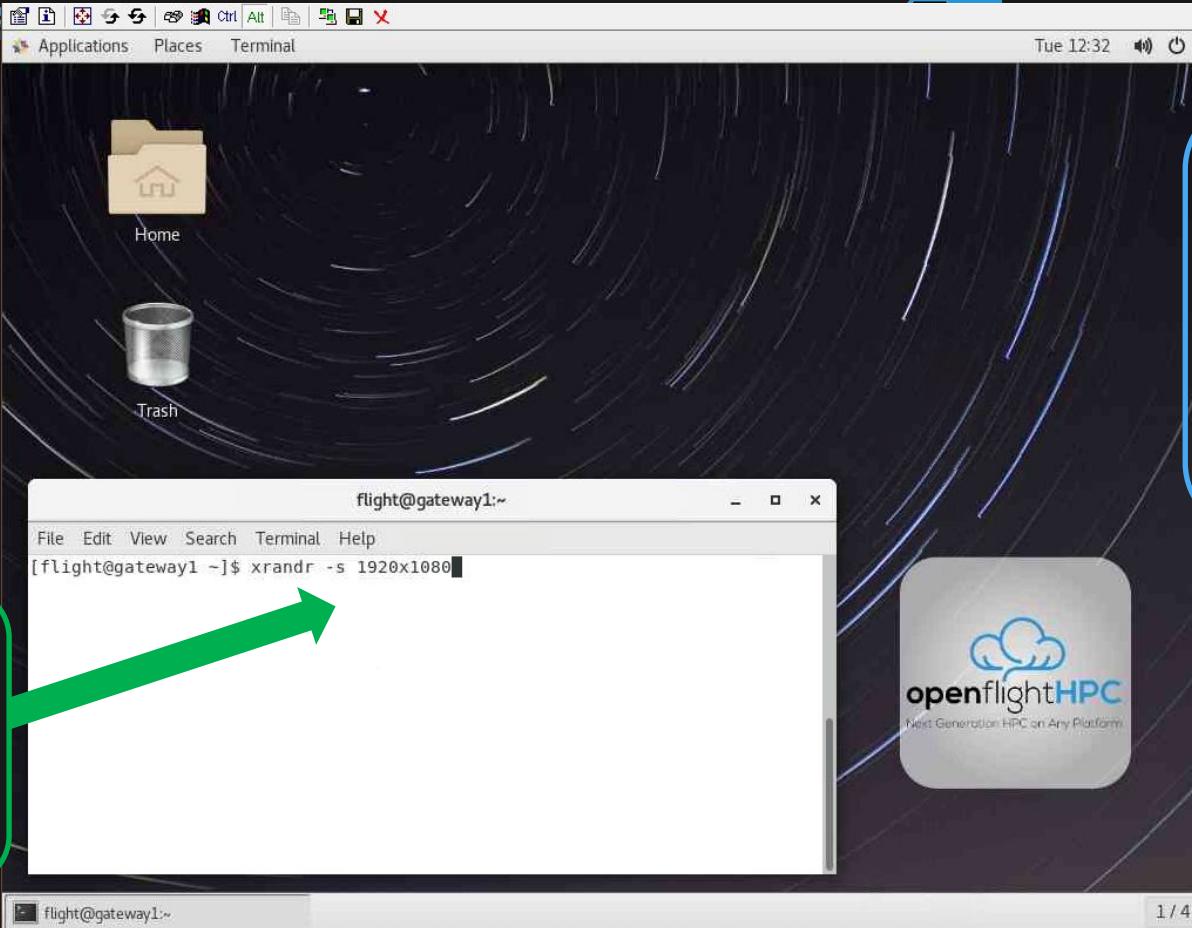
```
Identity: df0181a3-a800-473c-81ae-809442f
Type: gnome
Host IP: 51.11.13.120
Hostname: gateway1
Port: 5901
Display: :1
Password: Kk7qLu6s
```



This desktop session is directly accessible from the public internet. Depending on your client and network configuration you may be able to directly connect to the session using:

```
vnc://flight:Kk7qLu6s@51.11.13.120:5901
51.11.13.120:5901
51.11.13.120:1
```

# Starting a graphical desktop



Select a  
higher  
resolution  
if desired

*Hardware  
accelerated  
GPU rendering  
also available  
with compatible  
platform*

# Getting help to install an application



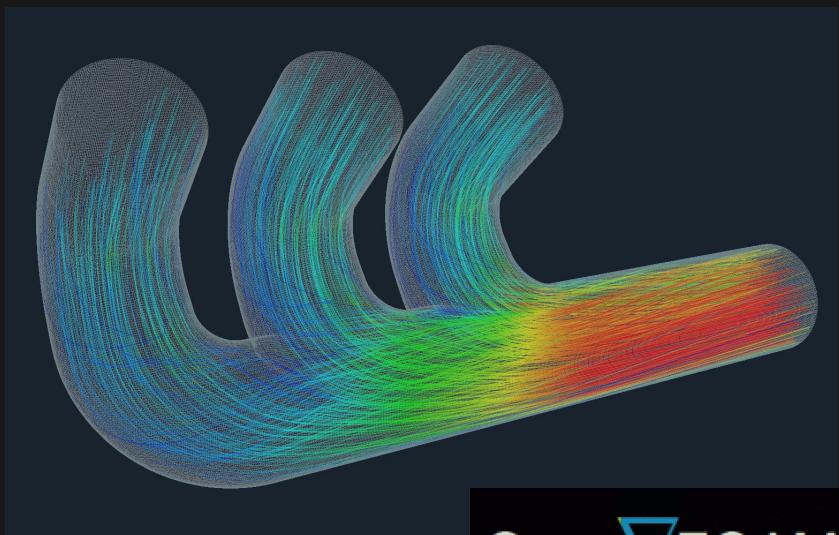
- Ask your cluster administrator
  - Typically for a hardware-based service (e.g. Barkla)
- For your personal HPC cluster
  - Use an application management environment
    - Conda
    - Easybuild
    - Gridware
    - Spack
  - Use a container service
    - Singularity (imports Docker modules)



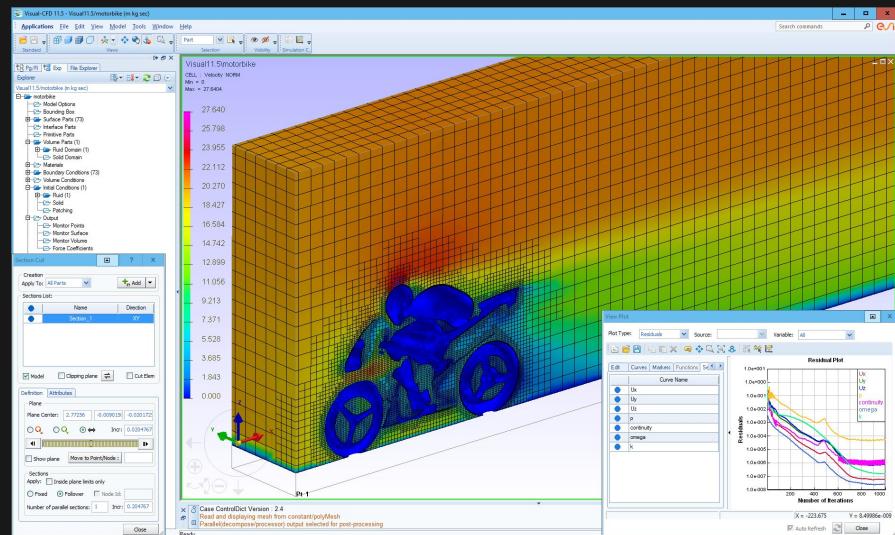
# Computational fluid dynamics with OpenFoam



- Popular engineering application
- Simulates fluid flow around objects



Open▽FOAM





- Distributed as source code
- Add-in toolboxes to support different functions
- Uses MPI for parallelism
- Most workflows are a multi-stage process
  - Interactive design stage to setup the problem
  - Decomposition of the problem into manageable pieces
  - Parallel processing of the decomposed mesh
  - Recombination of the parts into a single result
  - Visualisation of the output
- Relatively complex to install

Open $\nabla$ FOAM



# Installing applications with Gridware

```
# flight env create gridware
# flight env activate gridware
# gridware search openfoam
# gridware install apps/openfoam/4.1
```

Search limited to  
binary install  
versions by default

Press "y" to  
install deps

- *Flight will automatically detect and install software dependencies using binary downloads where possible*
- *Software is installed centrally and available on all nodes*
- *Compilation on demand is also possible if required*
- *Wait a few minutes for this stage to complete...*

Why?

# Running an OpenFoam job

```
#!/bin/bash
#SBATCH -N 1
flight env activate gridware
module load apps/openfoam
cp -r $FOAM_TUTORIALS/incompressible/icoFoam/cavity/cavity $HOME/.
cd $HOME/cavity
blockMesh
checkMesh
icoFoam
```

```
# module load apps/openfoam
# icoFoam –help
# nano myfoamjob.sh
```

Enter this content

Press **CTRL+X** to exit and enter “y” to save

Or – if that’s too much typing:

```
# curl -L http://tiny.cc/foamjob > myfoamjob.sh
```

# Submitting the OpenFoam job to run

```
# sbatch myfoamjob.sh  
# squeue
```

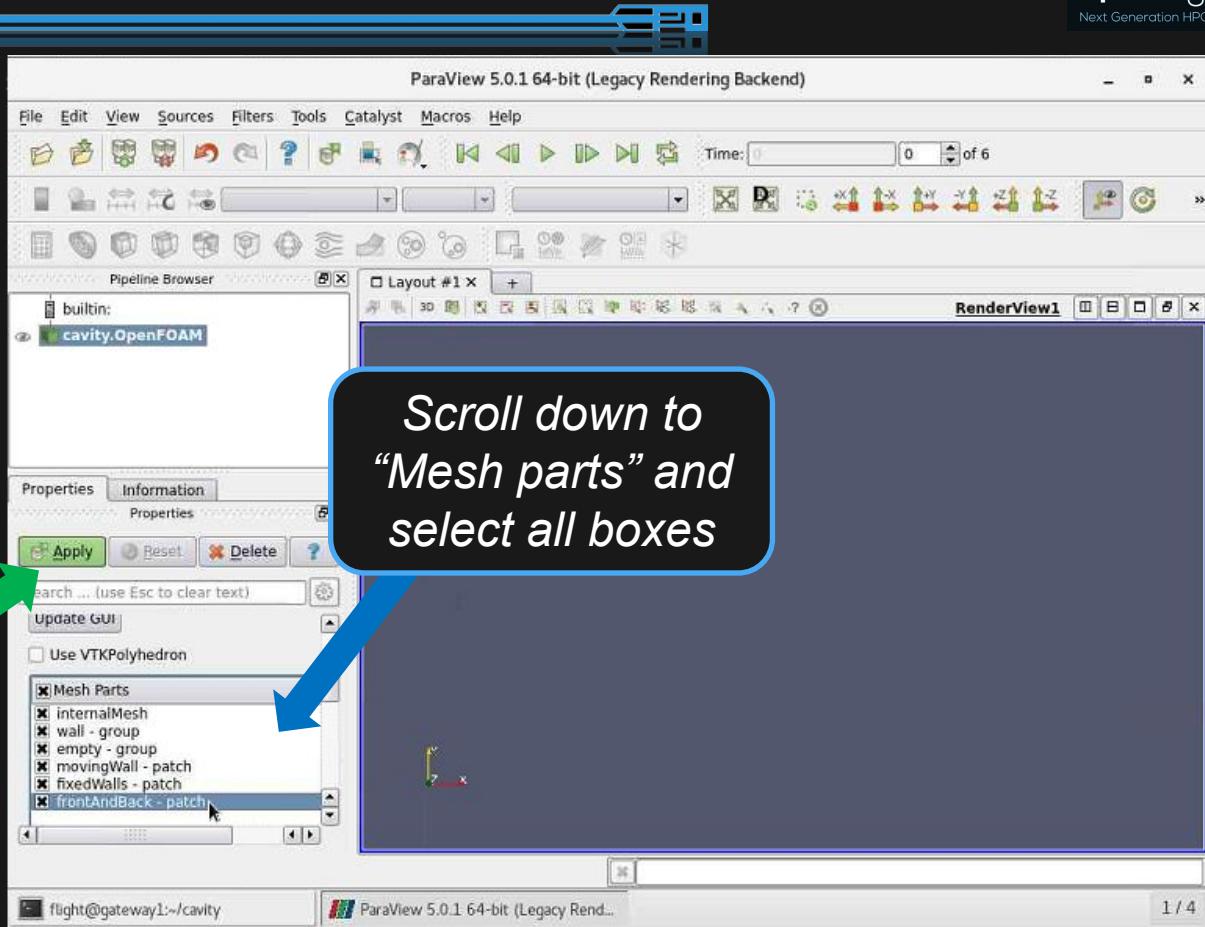
*Job should execute in 1-2 minutes  
on cluster compute nodes*

*Start a new  
desktop terminal  
and launch GUI*

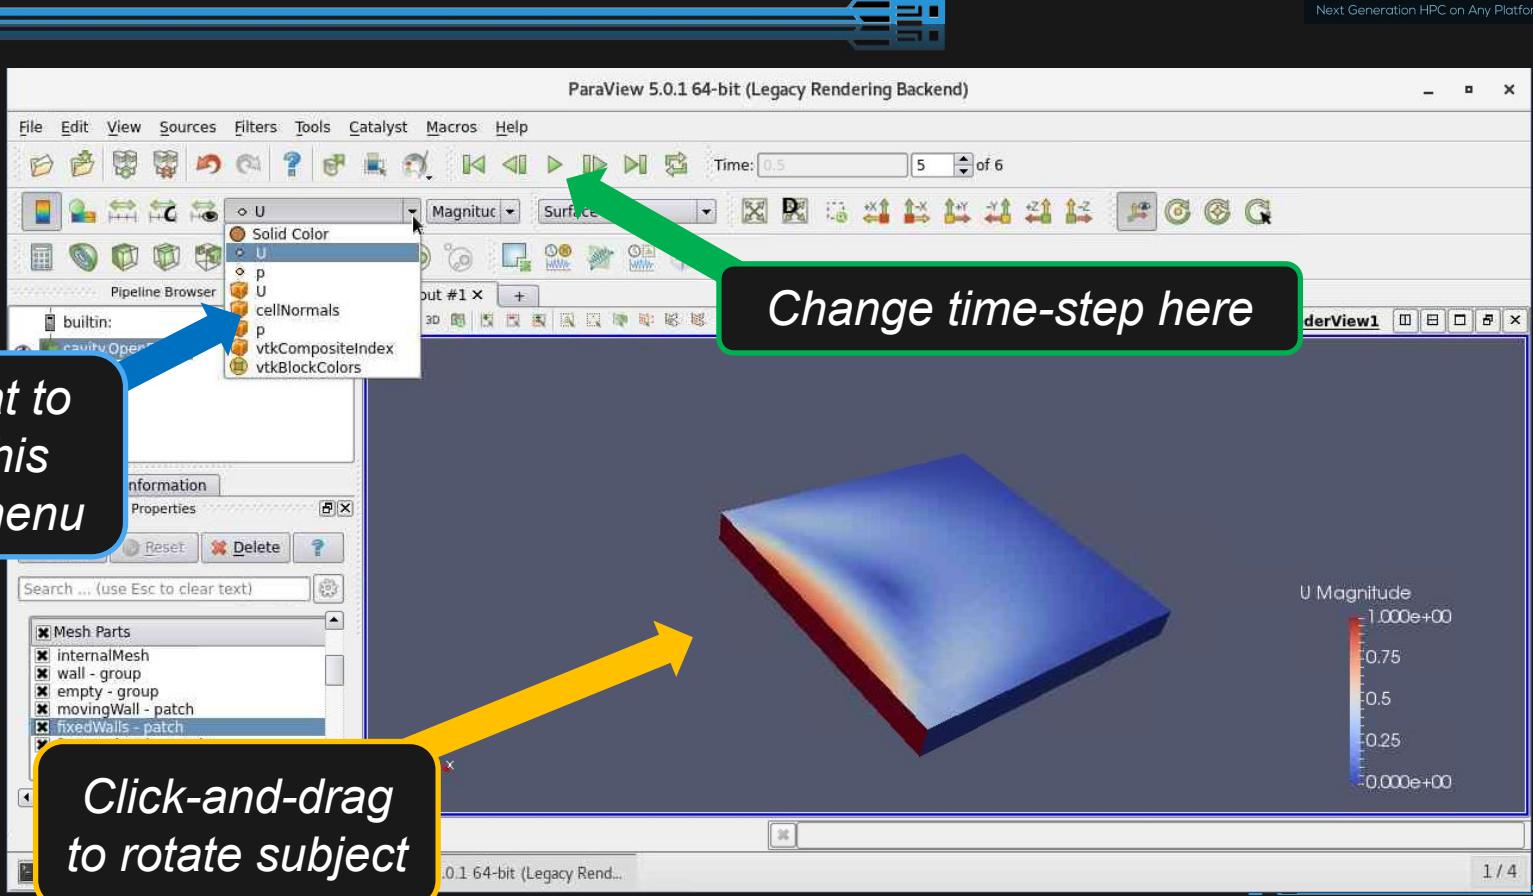


```
flight@gateway1:~  
File Edit View Search Terminal Help  
[flight@gateway1 ~]$ flight env activate gridware  
<gridware> [flight@gateway1 ~]$ module load apps/openfoam  
<gridware> [flight@gateway1 ~]$ cd cavity  
<gridware> [flight@gateway1 cavity]$ paraFoam
```

# Navigating the GUI



# CFD visualization



# Parallel computing jobs on an HPC cluster



- A parallel job uses resources from multiple nodes
  - CPU cores
  - Shared or distributed memory
  - Special devices (e.g. GPU, FPGA, etc.)
- Most parallel jobs use an API
  - Message-passing interface (MPI)
  - Provides simple routines for running code in parallel
- Compile your application with relevant libraries
  - MPIs: OpenMPI, Intel MPI, MPICH, MVAPICH
  - Threading: OpenMP, Intel TBB, SHMEM

# Writing some parallel C code



```
#include <stdio.h>
#include <mpi.h>
#include <time.h>
#include <string.h>

int main(int argc, char **argv) {
    char name[MPI_MAX_PROCESSOR_NAME];
    int nprocs, procno, len;
    MPI_Init( &argc, &argv );
    MPI_Comm_size( MPI_COMM_WORLD, &nprocs );
    MPI_Comm_rank( MPI_COMM_WORLD, &procno );
    MPI_Get_processor_name( name, &len );
    name[len] = '\0';
    time_t lt;
    lt = time(NULL);
    printf( "Hello !! from %s@%d/%d on %s\n", name, procno, nprocs, ctime(&lt));
    MPI_Barrier( MPI_COMM_WORLD );
    MPI_Finalize();
    return( 0 );
}
```

```
# module list
# which mpirun
# nano mympicode.c
```

Enter this content

*Or – if that's too much typing:*

```
# curl -L http://tiny.cc/mpienterexample > mympicode.c
```

# Compile your parallel code and submit

```
# module list
# which mpicc
# mpicc -o mympiapp -O3 mympicode.c
# ldd mympiapp
# mpirun -np 4 ./mympiapp
# nano mympijob.sh
```

Note: this is a capital "o" for optimisation (not zero)

```
# sbatch mympijob.sh
Submitted batch job 4
# cat slurm-4.out
```

Enter this content

Use this number here

```
#!/bin/bash
#SBATCH -N 2
flight env activate gridware
module load
mpi/openmpi/1.10.2
mpirun ./mympiapp
```

Press CTRL+X to exit  
and enter "y" to save

# More application environments

```
# flight env deactivate  
# flight env create conda  
# flight env activate conda  
# conda install perl  
# which perl  
# nano myperlscript.pl
```

Press “y”  
to confirm

Enter this  
content

```
use strict;  
use warnings;  
print "hi NAME\n";
```

*Press CTRL+X to exit and  
enter “y” to save*

# Conda application environment

```
# perl myperlscript.pl  
  
# nano myperljob.sh
```

Enter this content

```
#!/bin/bash -l  
#SBATCH -N 1  
echo Hello from $HOSTNAME  
flight env activate conda  
which perl  
perl myperlscript.pl
```

Press *CTRL+X* to exit and enter “y” to save

```
# sbatch myperljob.sh  
Submitted batch job 5  
  
# cat slurm-5.out
```

Use this number here

# Installing and running a web-service



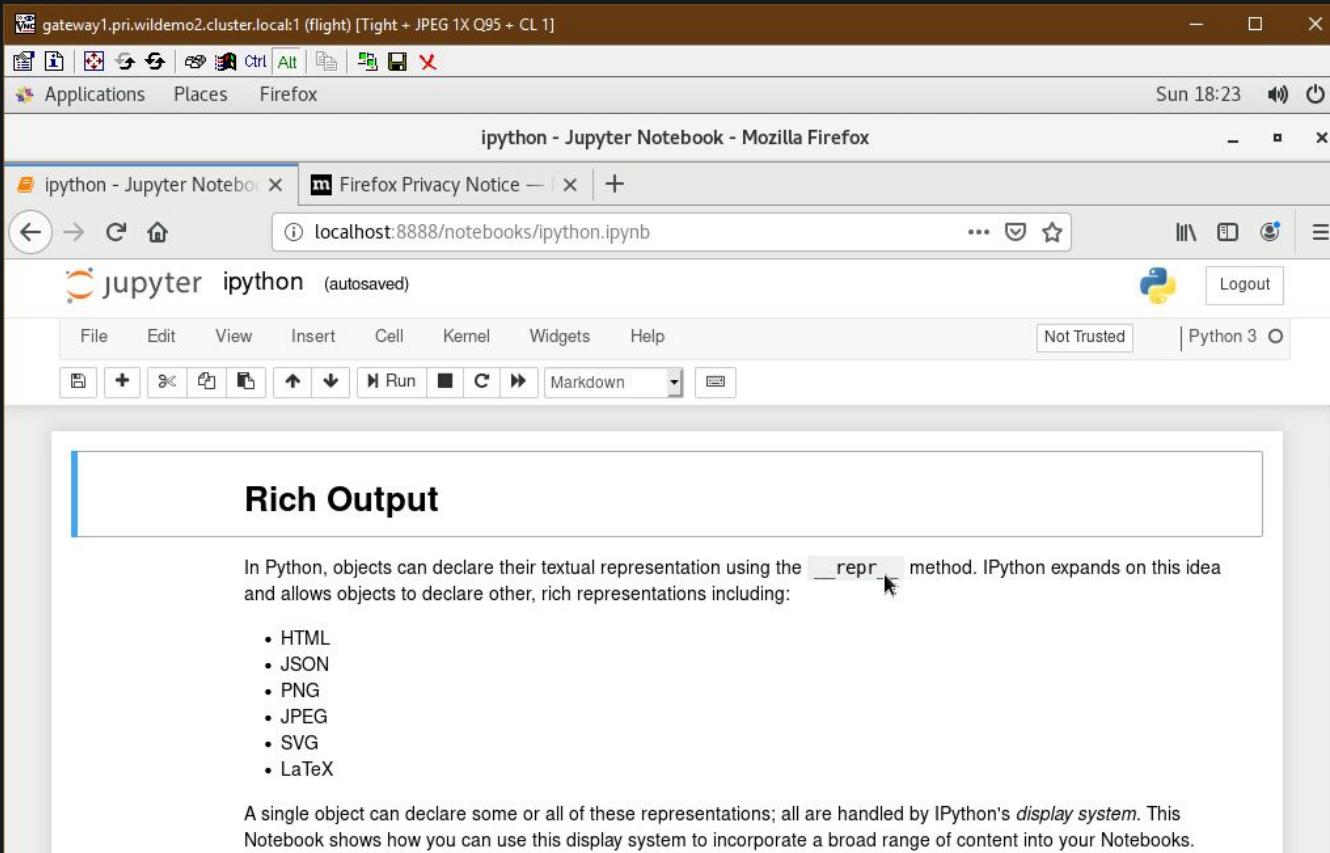
*Start a new  
desktop terminal  
and enter:*

File Edit View Search Terminal Help

flight@gateway1:~

```
[flight@gateway1 ~]$ flight env deactivate
[flight@gateway1 ~]$ flight env activate conda
(base) <conda> [flight@gateway1 ~]$ conda install jupyter
(base) <conda> [flight@gateway1 ~]$ jupyter --version
jupyter core      : 4.6.1
jupyter-notebook : 6.0.3
qtconsole        : 4.6.0
ipython          : 7.12.0
ipykernel         : 5.1.4
jupyter client   : 5.3.4
jupyter lab       : not installed
nbconvert         : 5.6.1
ipywidgets        : 7.5.1
nbformat          : 5.0.4
traitlets         : 4.3.3
(base) <conda> [flight@gateway1 ~]$ curl -L http://tiny.cc/jupyterexample > ipython.ipynb
(base) <conda> [flight@gateway1 ~]$ jupyter notebook ipython.ipynb
```

# Jupyter Notebook browser session



The screenshot shows a Jupyter Notebook interface running in Mozilla Firefox on a Linux desktop. The window title is "ipython - Jupyter Notebook - Mozilla Firefox". The address bar shows "localhost:8888/notebooks/ipython.ipynb". The notebook interface has a toolbar with icons for file operations, cell types (Code, Markdown, etc.), and execution. A menu bar includes File, Edit, View, Insert, Cell, Kernel, Widgets, and Help. A status bar at the bottom indicates "Not Trusted" and "Python 3". The main content area displays a section titled "Rich Output" with the following text:

In Python, objects can declare their textual representation using the `__repr__` method. IPython expands on this idea and allows objects to declare other, rich representations including:

- HTML
- JSON
- PNG
- JPEG
- SVG
- LaTeX

A single object can declare some or all of these representations; all are handled by IPython's *display system*. This Notebook shows how you can use this display system to incorporate a broad range of content into your Notebooks.

# Singularity for Docker containers



```
# flight env deactivate  
# flight env create singularity  
# flight env activate singularity  
# singularity run library://syllabsed/examples/lolcow
```

Close your  
*paraFoam* app  
before running

```
flight@gateway1:~> <singularity> [flight@gateway1 (wildemol) ~]$ singularity run library://syllabsed/examples/lolcow  
INFO: Convert SIF file to sandbox...  
WARNING: underlay of /etc/localtime required more than 50 (76) bind mounts  
  
/ Fine day to work off excess energy. \/  
\ Steal something heavy.  
-----  
 \ ^ ^  
  (oo)\_____  
   (__)\ )\/\|  
    ||----w |  
    ||     ||  
INFO: Cleaning up image...  
<singularity> [flight@gateway1 (wildemol) ~]$ █
```

# Running singularity jobs

```
# singularity pull shub://singularityhub/hello-world  
# nano mydockerjob.sh
```

Enter this content

```
#!/bin/bash -l  
#SBATCH -N 1 --mem=0  
flight start  
flight env activate singularity  
singularity run hello-world_latest.sif
```

```
# sbatch mydockerjob.sh
```

Submitted batch job 6

```
# cat slurm-6.out
```

Press *CTRL+X* to exit and enter “y” to save

Use this number here

# Spack application management



```
# flight env deactivate  
# flight env create spack  
# flight env activate spack  
# spack list  
# spack install bowtie  
# nano mybowtiejob.sh
```

*Note: this is a capital “o” for output (not zero)*

Enter this content

```
#!/bin/bash -l  
#SBATCH -N 1  
flight env activate spack  
spack load bowtie  
wget -O ecoli.fq http://tiny.cc/ecoli  
bowtie-build ecoli.fq e_coli
```

```
# sbatch mybowtiejob.sh
```

Submitted batch job 7

```
# cat slurm-7.out
```

Press **CTRL+X** to exit and enter “y” to save

Use this number here

# Easybuild application management



```
# flight env deactivate  
# flight env create easybuild  
# flight env activate easybuild  
# module load EasyBuild  
# eb -S R  
# eb R-3.5.1-foss-2018b-Python-2.7.15.eb -r
```

*Robot mode – automatically resolve and install dependencies*

*Easybuild may take some time to complete the build, as it will pull in any dependencies needed*

*Once completed, applications will be available on all compute nodes of the cluster*

# Review



So far we have:

- Created our own applications from source
- Used various application environments to install apps
- Run jobs and visualized data using a compute cluster

What if we want to run:

- A MapReduce workflow?
- A machine-learning workflow?

# Installing Hadoop

Note: this is a capital “o” for output (not zero)

```
# flight env deactivate
# sudo yum install java-1.8.0-openjdk.x86_64
# sudo yum install java-1.8.0-openjdk-devel.x86_64
# wget -O /tmp/hadoop.tgz http://tiny.cc/hadoop321
# cd /opt/apps
# tar xzf /tmp/hadoop.tgz
# cd hadoop-3.2.1
```

Press “y”  
to confirm

# Configuring Hadoop

```
# nano etc/hadoop/hadoop-env.sh
```

*Press **CTRL+U** (underscore) and enter 54 to jump to this line*

*Edit line 54 so it reads:*

```
export JAVA_HOME=/usr/lib/jvm/java-1.8.0-openjdk-1.8.0.242.b08-0.el7_7.x86_64/jre
```

*Press **CTRL+X** to exit and enter “y” to save*

*Or – if that’s too much typing:*

```
# curl -L https://tinyurl.com/hdconfig | bash
```

# Preparing our work

Note: this is a capital “o” for output (not zero)

```
# cd  
# wget https://tinyurl.com/hadoopenv  
# source hadoopenv  
# mkdir MapReduceTutorial  
# chmod 777 MapReduceTutorial  
# cd MapReduceTutorial  
# wget -O hdffiles.zip https://tinyurl.com/hdinput1  
# unzip -j hdffiles.zip  
# ls
```

*This archive includes a spreadsheet of data which we will sort into sales units per region*

# Compiling java + setting manifest

```
# pwd  
/home/flight/MapReduceTutorial  
# javac -d . SalesMapper.java SalesCountryReducer.java  
SalesCountryDriver.java  
# nano Manifest.txt
```

*This should be a single long line*

Enter this content

Main-Class: SalesCountry.SalesCountryDriver

*Press CTRL+X to exit and enter “y” to save*

# Create the jar file



```
# jar cfm ProductSalePerCountry.jar Manifest.txt  
SalesCountry/*.class  
# ls *.jar  
ProductSalePerCountry.jar
```

*This should be a  
single long line*

This is our java archive file  
containing the code ready to run

# Start Hadoop and copy in data-files

```
# $HADOOP_HOME/sbin/start-dfs.sh  
# $HADOOP_HOME/sbin/start-yarn.sh
```

*Distributed file system service*

*Resource manager, node manager, app manager*

*Copy in your data-file to the DFS*

```
# mkdir ~/inputMapReduce  
# cp SalesJan2009.csv ~/inputMapReduce/.  
# $HADOOP_HOME/bin/hdfs dfs -ls ~/inputMapReduce
```

# Run your JAR through Hadoop

*Execute your MapReduce job*

```
# $HADOOP_HOME/bin/hadoop jar  
ProductSalePerCountry.jar ~/inputMapReduce  
~/mapreduce_output_sales
```

*This should be a single long line*

*Check that output was produced*

```
# $HADOOP_HOME/bin/hdfs dfs -cat  
~/mapreduce_output_sales/part-00000 | more
```

*This should be a single long line*

# Running an AI workload in Tensorflow



```
# flight env deactivate  
# flight env activate singularity  
# git clone https://github.com/tensorflow/models.git  
# singularity exec docker://tensorflow/tensorflow:1.15.0  
python ./models/tutorials/image/mnist/convolutional.py
```

*This should be a single long line*

*Version 1.15.0 of Tensorflow runs on CPU by default*

*This example model will take 5-10 minutes to prepare and run for 20-30 minutes*



# Alternatively, build in Conda



```
# flight env deactivate  
# flight env activate conda  
# conda create -n tensorflow python=3.6  
# source activate tensorflow  
# pip install tensorflow==1.15
```

*Version  
1.15.0 of  
Tensorflow  
runs on CPU  
by default*

```
# git clone https://github.com/tensorflow/models.git  
# python ./models/tutorials/image/mnist/convolutional.py
```

# Summary

## The Hybrid HPC Experience:

- Worked in a fixed cluster environment.
  - Template designed for training.
  - Easy to replicate, easy to mimic HPC clusters such as Barkla.
- Engaged with a special HPC use case using Hadoop.
  - Aimed at project-based HPC work.
  - Designed to deal with unusual or periodic requests.

# Hybrid HPC leverages the right resources at the right time.

# What next?



- Your cluster is active until:
- Have a project? Idea? Concept? Talk to us!

Cristin Merritt: [Cristin.Merritt@alces-flight.com](mailto:Cristin.Merritt@alces-flight.com)

Wil Mayers: [Wil.Mayers@alces-flight.com](mailto:Wil.Mayers@alces-flight.com)



# openflightHPC

Next Generation HPC on Any Platform



alcesflight



© Alces Flight Ltd 2020– Page 93